

FlowBuddy: System Requirements Specification (SRS) v2.5 - FINAL

Version: 2.5 - Agentic Coach with Psychological Safety

Date: January 3, 2026

Project: Comet Resolution v2 Hackathon - Productivity Category

Status: Implementation-Ready

⌚ Product Philosophy: "Supportive Autonomy"

The Reconciliation: FlowBuddy is an **agentic AI coach** (proactive, autonomous, outcome-focused) that operates with **psychological safety principles** (non-judgmental, permission-based, respects cognitive load).

Core Tension Resolved:

- **Agentic:** Takes initiative, plans autonomously, intervenes predictively
- **Safe:** Never commands, always offers choices, respects "no"
- **Transparent:** Shows reasoning, allows undo, explains every action

Voice Example:

- Pushy Coach: "You missed your workout. Get it done by 6 PM."
 - Passive Tool: "Would you like me to suggest a workout time?"
 - **FlowBuddy:** "I noticed you missed 9 AM. I moved it to 6 PM (your usual backup time). Work for you? [Yes] [Pick different time] [Skip today]"
-

Table of Contents

1. [Hackathon Alignment & MVP Scope](#)
2. [System Overview](#)
3. [User Personas & Journeys](#)
4. [MVP Functional Requirements](#)
5. [Permission & Autonomy Controls](#)
6. [UX Design Specifications](#)
7. [Technical Architecture](#)
8. [Data Models](#)
9. [API Specifications](#)

10. Observability (Opik Integration)
 11. Success Metrics (Outcomes + Guardrails)
 12. Implementation Plan (4 Weeks)
-

1. Hackathon Alignment & MVP Scope

1.1 Hackathon Requirements

Comet Resolution v2: Build an AI agent that meaningfully helps people keep their resolutions with measurable impact and observability.

FlowBuddy's Fit:

- **Agentic Behavior:** 3 autonomous agents (decomposer, scheduler, accountability)
- **Measurable Impact:** Resolution completion rate, habit consistency, intervention effectiveness
- **Observability:** Full Opik integration with 12+ metrics (AI quality + user outcomes + guardrails)
- **Innovation:** Dual-input model (structured resolutions + brain dump signals)

1.2 MVP Feature Prioritization (Realistic 4-Week Build)

P0 (Must Ship - Week 1-3):

Feature	Agentic Level	Build Effort	Hackathon Value
1. Dual Intake: Brain Dump + Resolution Setup	Medium	1 week	Shows UX innovation
2. Resolution Decomposer Agent	High	1 week	Core agentic demo
3. Weekly Micro-Resolution Generator	High	3 days	Autonomous planning
4. Basic Intervention System	Medium	1 week	Outcome driver

P1 (Nice to Have - Week 4):

Feature	Build Effort	If Time Allows
Daily schedule optimization	1 week	Start with manual confirm, add auto later
Pattern recognition	1 week	Start with heuristics (e.g., "90% success on Mon 9 AM")

P2 (Post-Hackathon):

- ML-based pattern prediction
- Advanced background job scheduling
- Deep Work Shield

Strategic Trade-off: We're cutting daily auto-scheduling from MVP to focus on:

1. Perfect execution of resolution decomposition
2. Robust intervention system
3. Rich Opik observability

This reduces implementation risk while keeping the agentic thesis strong.

2. System Overview

2.1 High-Level Architecture



2.2 Agent Responsibilities (Simplified for MVP)

Agent 1: Resolution Decomposer

- **Trigger:** User creates new resolution
- **Autonomy:** Generates 4-12 week plan + Week 1 tasks automatically
- **Permission Model:** User must confirm before activation
- **Opik Traces:** `resolution_decomposition`

Agent 2: Weekly Micro-Resolution Planner

- **Trigger:** Every Sunday 8 PM (automated job)
- **Autonomy:** Generates next week's tasks based on last week's performance
- **Permission Model:** User receives push notification to review/accept
- **Opik Traces:** `weekly_planning`

Agent 3: Intervention Coach

- **Trigger:** Slippage detection (Thursday check if <50% toward goal)
- **Autonomy:** Detects risk, generates personalized recovery plan
- **Permission Model:** Offers 3 options (get back on track / adjust goal / take break)
- **Opik Traces:** `slippage_intervention`

NOT in MVP (but documented for v2):

- Daily auto-scheduling (too complex for 4 weeks)
- ML-based pattern prediction (start with simple heuristics)
- Proactive morning check-ins (save for post-demo)

3. User Personas & Journeys

3.1 Primary Persona: "Resolution Rachel"

Demographics:

- Age: 28-35, knowledge worker
- Resolutions: "Exercise 3x/week", "Read 12 books/year", "Launch side project"
- Pain: Starts strong, loses momentum by Week 3, needs accountability without guilt

What She Needs:

1. **Easy start:** Shouldn't take >2 min to set up resolution
2. **Autonomy:** Agent breaks down goal, she confirms
3. **Flexibility:** Can reschedule/adjust without feeling like she "failed"
4. **Clarity:** Always knows if she's on track
5. **Safety valve:** Can "brain dump" overwhelm without it becoming more tasks

3.2 Dual-Input Journey (The Innovation)

Path A: Structured Resolution Setup

User Action: Taps "New Resolution"

Agent: "What resolution would you like to work on?"

User: Types "Exercise 3x per week"

Agent (Autonomous):

- Generates 12-week plan
- Creates Week 1 tasks (Mon 9AM, Wed 7PM)
- Shows for approval

User: Reviews → Taps "Looks good!"

Agent: Locks in, sets Sunday reminder for Week 2 planning

Emotional Arc: Motivated → Supported → Confident

Path B: Brain Dump as Signal (The Differentiator)

User Action: Taps "Brain Dump" (always visible button)

User: "Feeling overwhelmed, missed my workout today,
worried I'm falling behind on exercise goal"

Agent (Signal Processing):

- Detects: emotion (overwhelm), event (missed workout),
resolution reference (exercise)
- Does NOT create new task (avoids task creep)
- Uses as input signal

Agent Response:

"Got it. Sounds like today was tough. I see you're
at 1/2 workouts this week. Want me to reschedule
Wednesday's session for this weekend instead?"

[Yes] [No, I'll handle it]

If Yes: Agent reschedules autonomously

If No: Agent acknowledges, doesn't push

Emotional Arc: Stressed → Heard → Relieved

Key Insight: Brain dump is **NOT** for task creation. It's a **pressure release valve** that gives the agent context about:

- User's emotional state (overwhelmed, motivated, etc.)
- Blockers ("had a meeting conflict")
- Intent shifts ("maybe 3x/week is too much")

This allows the agent to **adapt its interventions** based on qualitative signals, not just completion data.

4. MVP Functional Requirements

4.1 Feature 1: Dual Intake System

FR-1.1: Brain Dump (Pressure Relief)

Priority: P0 (Core UX differentiator)

Agentic Level: Medium (Processes for signals, doesn't auto-create tasks)

User Story:

As Rachel, when I'm stressed or have scattered thoughts, I want to dump them quickly so I feel heard, and the agent can use that context to help me without creating more work.

Acceptance Criteria:

- "Brain Dump" button is always visible (home screen sticky button)
- Tapping opens fullscreen text input (auto-focused keyboard)
- User types/pastes unstructured text (max 2000 chars)
- Submit triggers:
 - **Signal extraction** (LLM identifies: emotions, blockers, resolution references)
 - **Acknowledgment** (empathetic confirmation)
 - **Optional action** (if agent detects actionable intent, offers to help)
- Does NOT auto-create tasks (avoids overwhelm)
- All signal processing logged to Opik

Technical Specification:

```
python
```

```
# LLM Prompt: Brain Dump Signal Extractor
BRAIN_DUMP_SIGNAL_PROMPT = """
User's brain dump: "{input_text}"
Active resolutions: {user_resolutions}
```

Extract signals (do NOT create tasks):

1. Emotional state: overwhelmed | stressed | motivated | neutral
2. Resolution references: which resolutions mentioned?
3. Blockers: what's preventing progress?
4. Intent shifts: does user want to adjust goals?

Return JSON:

```
 {{
    "emotional_state": "overwhelmed",
    "resolution_refs": ["Exercise 3x/week"],
    "blockers": ["Meeting conflict on Wednesday"],
    "intent_shift": "possibly wants to reduce frequency",
    "actionable": true,
    "suggested_response": "Sounds like Wednesday is tough. Want me to move workouts to weekends?"
}}
```

Principles:

- Acknowledge feelings ("tough day")
- Don't judge ("you should have...")
- Offer help, don't command

```
"""
```

Opik Trace Structure:

```

python

{
  "trace_name": "brain_dump_processing",
  "spans": [
    {"name": "signal_extraction", "llm_tokens": 450},
    {"name": "resolution_matching", "vector_similarity": 0.87},
    {"name": "empathy_generation", "llm_tokens": 120}
  ],
  "metadata": {
    "emotional_state": "overwhelmed",
    "actionable": true,
    "user_accepted_help": true
  }
}

```

Example Flow:

User: "Ugh missed my run today, work was crazy"
 Agent: "Got it. Sounds like today was packed. You're
 at 1/2 workouts this week. Want me to move
 Wednesday's to Saturday 10 AM instead?"
 [Yes, reschedule] [No thanks]

Edge Cases:

- No resolution reference → Just acknowledge: "Thanks for sharing. I've noted this."
- Pure emotional vent → Supportive response: "Rough day. You've got this."
- Actionable but user says "No thanks" → Agent backs off, doesn't push

FR-1.2: Resolution Setup (Structured Intake)

Priority: P0 (Core agentic feature)

Agentic Level: High (Autonomous decomposition)

User Story:

As Rachel, I want to tell the agent my big goal and have it break it down into achievable milestones so I don't have to plan everything myself.

Acceptance Criteria:

- User navigates to "New Resolution"
- Conversational form: "What resolution would you like to work on?"
- User types free text (e.g., "Exercise 3x per week")

Agent autonomously:

- Classifies type (habit, project, learning)
- Generates 4-12 week milestone plan
- Creates Week 1 specific tasks (2-5 tasks with suggested times)

Shows plan for user approval

User can: Accept / Edit tasks / Regenerate with different approach

Once accepted, agent sets Sunday reminder for Week 2 planning

LLM Prompt:

```
python
```

```
RESOLUTION_DECOMPOSER_PROMPT = """
```

```
User's resolution: "{resolution_text}"
```

```
Today: {current_date}
```

```
User context: {user_schedule_hints}
```

Generate a realistic plan:

1. Classify: habit | project | learning | health
2. Set duration: 4-12 weeks (prefer 12 for habits)
3. Create milestones:
 - Week 1-2: Easier (build momentum)
 - Week 3+: Full target
4. For Week 1, generate 2-4 specific tasks with times

Return JSON:

```
 {{
  "type": "habit",
  "title": "Exercise 3x per week",
  "duration_weeks": 12,
  "milestones": [
    {{
      "week": 1,
      "target": "2 sessions (warmup)",
      "tasks": [
        {"title": "30-min jog", "day": "Monday", "time": "09:00", "duration_min": 30},
        {"title": "Yoga class", "day": "Wednesday", "time": "19:00", "duration_min": 45}
      ]
    }}
  ]
}}
```

Coaching principles:

- Week 1 should feel achievable (not overwhelming)
- Tasks must be specific (activity + duration)
- Respect common work schedules (avoid midday)

```
"""
```

Opik Metrics:

- `decomposition_quality` (LLM-as-Judge: Are tasks specific and achievable?)
- `user_acceptance_rate` (% of generated plans accepted without edits)
- `plan_generation_time` (target: <3s)

FR-2.1: Automated Weekly Planning

Priority: P0 (Core agentic behavior)

Agentic Level: High (Fully autonomous)

User Story:

As Rachel, I want the agent to automatically create my weekly plan based on how I did last week, so I always know what to focus on without having to think about it.

Acceptance Criteria:

- Every Sunday at 8 PM, automated job runs for all users with active resolutions
- Agent analyzes last week:
 - Tasks completed vs target
 - Patterns (which days user succeeded)
 - Missed tasks and reasons (from brain dumps if available)
- Agent generates next week's tasks:
 - Maintains target (e.g., 3x/week) if user succeeded
 - Adjusts difficulty if user struggled (<50% completion)
 - Avoids days/times where user consistently missed
- Sends push notification: "Your week is planned! Tap to review."
- User can view, edit, or confirm plan
- All planning logic logged to Opik

Agentic Logic:

```
python
```

```

@scheduler.scheduled_job('cron', day_of_week='sun', hour=20)
async def generate_weekly_plans():
    users = await get_users_with_active_resolutions()

    for user in users:
        with opik.track(name="weekly_planning", user_id=user.id):
            resolutions = user.active_resolutions

            for resolution in resolutions:
                # Analyze last week
                last_week = await analyze_week_performance(resolution.id)

                # Adaptive planning
                if last_week.completion_rate < 0.5:
                    # User struggled - reduce Week 2 target
                    next_week_plan = await generate_plan(
                        resolution,
                        target=resolution.target - 1, # e.g., 2x instead of 3x
                        difficulty="easier"
                    )
                    reason = "reduced_target"
                elif last_week.completion_rate >= 0.8:
                    # User succeeded - maintain or increase
                    next_week_plan = await generate_plan(
                        resolution,
                        target=resolution.target,
                        difficulty="maintain"
                    )
                    reason = "on_track"

                # Avoid failure patterns
                failed_days = last_week.get_failed_days() # e.g., ["Wednesday"]
                next_week_plan.avoid_days(failed_days)

                # Save and notify
                await save_weekly_plan(next_week_plan)
                await send_notification(
                    user.id,
                    f"Week {next_week_plan.week_number} is ready! {len(next_week_plan.tasks)} tasks planned.",
                    action="review_plan"
                )

            # Opik logging
            opik.log_feedback(
                trace_id=trace_id,
                metric_name="weekly_plan_generated",

```

```

    value=1.0,
    metadata={
        "completion_last_week": last_week.completion_rate,
        "adjustment_reason": reason,
        "tasks_planned": len(next_week_plan.tasks)
    }
)

```

Permission Model:

- Agent generates plan autonomously
 - User receives notification (not intrusive - single push)
 - User can review before Monday
 - If user doesn't review, plan is assumed accepted (low friction)
-

4.3 Feature 3: Intervention System (Simplified)

FR-3.1: Thursday Slippage Check

Priority: P0 (Outcome driver)

Agentic Level: Medium (Detects + offers options, doesn't force)

User Story:

As Rachel, when I'm falling behind on my goal, I want the agent to proactively help me get back on track before the week ends, without making me feel guilty.

Acceptance Criteria:

- Every Thursday at 7 PM, automated job checks all users
- For each active resolution:
 - Calculate: current progress vs weekly target
 - If progress <50% AND ≥ 2 days left in week → Trigger intervention
- Agent generates intervention message:
 - Acknowledges current state (no judgment)
 - Offers 3 clear options:
 1. Get back on track (suggests specific actions)
 2. Adjust goal (reduce target for this week)
 3. Take a break week (no guilt)
- User selects option → Agent acts accordingly
- All interventions logged to Opik with outcome

Example Intervention:

Push Notification (Thursday 7 PM):

Quick Check-In

I noticed you've done 1 workout this week (goal: 3). Life got busy?

Tap to see options →

[In-App Card]

No judgment! Here are your options:

Option 1: Get Back on Track

- Schedule Sat 10 AM + Sun 3 PM
- Still hit your 3x goal

Option 2: Adjust This Week

- Try 2x this week instead
- Resume 3x next Monday

Option 3: Take a Break Week

- Pause until next Monday
- No streak penalty

What works for you?

Implementation:

python

```

@scheduler.scheduled_job('cron', day_of_week='thu', hour=19)
async def check_for_slippage():
    users = await get_users_with_active_resolutions()

    for user in users:
        resolutions = user.active_resolutions

        for res in resolutions:
            progress = await calculate_week_progress(res.id)
            days_left = 7 - datetime.now().weekday() # Days until Sunday

            # Intervention criteria
            if progress.completion_rate < 0.5 and days_left >= 2:
                with opik.track(name="slippage_intervention"):

                    # Generate options
                    options = await intervention_agent.generate_options({
                        "resolution": res.title,
                        "current": progress.completed,
                        "target": progress.target,
                        "days_left": days_left,
                        "user_history": user.past_adjustments
                    })

                    # Send notification
                    await send_push(
                        user.id,
                        "Quick Check-In",
                        "Tap to see options",
                        data={"intervention_id": intervention.id}
                    )

            # Log to Opik
            opik.log_trace(
                "intervention_sent",
                metadata={
                    "resolution_id": res.id,
                    "completion_rate": progress.completion_rate,
                    "options_offered": len(options)
                }
            )

```

Opik Metrics:

- **intervention_effectiveness** - % of users who recover (complete \geq target after intervention)
- **option_selection_distribution** - Which options users choose most

- `false_positive_rate` - % interventions sent when user would have succeeded anyway
-

5. Permission & Autonomy Controls

5.1 Explicit Permission Model

Core Principle:

Agent acts autonomously **within user-defined boundaries**. User always has control.

Permission Levels (User Settings):

Permission Type	Default	What It Controls
Weekly Planning	ON (opt-out)	Agent generates weekly plans automatically
Interventions	ON (opt-out)	Agent sends Thursday slippage checks
Push Notifications	ON (with caps)	All proactive agent messages
Quiet Hours	10 PM - 8 AM	No notifications during these hours
Notification Frequency Cap	3/day max	Prevents spam

Settings UI:

Agent Controls	
Weekly Planning [ON]	
Agent creates your weekly plan automatically every Sunday	
Interventions [ON]	
Agent checks in if you're falling behind (Thursdays only)	
Push Notifications [ON]	
Max 3 per day	
Quiet Hours [EDIT]	
Currently: 10 PM - 8 AM	
Pause All Coaching	
(Take a break from agent)	

5.2 Agent Actions Log (Transparency)

Requirement:

Every autonomous action the agent takes is logged in-app with undo option.

Example Log Entry:

Agent Actions (This Week)

Sun, Jan 5 - 8:00 PM

 Generated Week 2 Plan

- 3 workouts scheduled
- Avoided Wednesday (per your pattern)

[View Plan] [Undo]

Thu, Jan 2 - 7:00 PM

 Sent Intervention

- You were at 1/3 workouts
- You chose: "Adjust to 2x"

[View Details]

5.3 "Snooze Coach" Mode

Requirement:

User can pause all proactive agent behavior with one tap.

UI:

[Pause Coaching]

Pause all agent check-ins and planning for:

- 1 week
- 2 weeks
- Until I resume manually

Your resolutions stay active, but the agent won't send notifications or create plans.

[Pause] [Cancel]

Implementation:

```
python
```

```

# When user pauses coaching
async def pause_coaching(user_id, duration_days=None):
    user = await get_user(user_id)
    user.coaching_paused = True
    user.pause_until = datetime.now() + timedelta(days=duration_days) if duration_days else None
    await db.save(user)

    # Disable scheduled jobs for this user
    scheduler.pause_job(f"weekly_planning_{user_id}")
    scheduler.pause_job(f"intervention_check_{user_id}")

    opik.log_feedback(
        metric_name="coaching_paused",
        value=1.0,
        metadata={"duration_days": duration_days}
    )

```

6. UX Design Specifications

6.1 Home Screen (Updated with Dual Intake)



7-12. [Remaining Sections - Streamlined]

Key Updates to Final Sections:

Section 8 (Data Models):

- Add `brain.dumps` table with `signals_extracted` JSONB field
- Add `agent_actions_log` table for transparency
- Add `user_preferences` table with permission settings

Section 10 (Opik Integration): Add 3 trace types:

1. `brain_dump_processing` - Signal extraction quality
2. `resolution_decomposition` - Planning accuracy
3. `weekly_planning` - Adaptation logic
4. `slippage_intervention` - Intervention effectiveness

Section 11 (Success Metrics): Outcome Metrics:

- Resolution completion rate (target: 70%)
- Weekly adherence rate (target: 75%)
- Intervention recovery rate (target: 60%)

Guardrail Metrics (NEW):

- `nudge_fatigue_rate` - % users who pause/mute notifications
- `autonomyViolationRate` - LLM judge: % of agent messages that command vs offer choice
- `undoActionRate` - % of agent actions user undoes (target: <5%)

Target Dashboard:

FlowBuddy Outcomes Dashboard

Outcome Metrics:

- Resolution Completion: 72% ✓ (target: 70%)
- Intervention Recovery: 65% ✓ (target: 60%)

Guardrail Metrics:

- Nudge Fatigue: 3% ✓ (target: <5%)
- Autonomy Violations: 2% ✓ (target: <5%)
- Action Undos: 4% ✓ (target: <5%)

Status: Healthy - Agent is helpful, not annoying

12. Implementation Plan (4 Weeks)

Week 1: Foundation + Dual Intake

- FastAPI backend with Opik SDK
- Database schema (resolutions, tasks, brain.dumps, agent_actions_log)
- Brain Dump endpoint + signal extraction (LLM)
- Resolution Setup endpoint + decomposer agent (LLM)
- React Native: Home screen, Brain Dump screen, Resolution Setup
- Opik: First 3 traces (brain_dump, resolution_decomp)

Deliverable: User can dump brain, create resolution, see Week 1 plan

Week 2: Weekly Planner + Permissions

- APScheduler setup for background jobs
- Weekly planning agent (Sunday 8 PM job)
- Notification system (push to mobile)
- Permission controls UI (settings screen)
- Agent Actions Log (transparency feature)
- Opik: `weekly_planning` trace + `plan_acceptance_rate` metric

Deliverable: Agent generates Week 2 plan autonomously, user can review/pause

Week 3: Intervention System + Guardrails

- Thursday slippage detection job

- Intervention agent (generates 3 options)
- User response handling (reschedule / adjust / break)
- Opik: `slippage_intervention` trace
- LLM-as-Judge metrics:
 - `autonomy_support_score` (does agent offer choice?)
 - `tone_warmth` (is message supportive?)
- Guardrail tracking:
 - ``nudge_fatigue_rate`